

data_aug_adam_val_95_28

October 7, 2021

1 Covid Classifier Model

1.0.1 Goals

Classify: - Normal CXR - Viral Pneumonia CXR - COVID CXR

1.1 Create Directories for Dataset

Separate the data to use later as generators.

```
[ ]: # Matriz de confusion, cambiar learnings rates (learning rates dinamicos),  
      ↳ dropouts 0.3 & 0.2, Batch Normalization  
# K-Fold o avg de modelos  
import os  
  
BASE_PATH = '/home/hivini/learn/research/new-covid'  
ORIGINAL_DATASET_DIR = os.path.join(BASE_PATH, 'COVID-19_Radiography_Dataset')  
ORIGINAL_VIRAL_DIR = os.path.join(ORIGINAL_DATASET_DIR, 'Viral Pneumonia')  
ORIGINAL_COVID_DIR = os.path.join(ORIGINAL_DATASET_DIR, 'COVID')  
ORIGINAL_NORMAL_DIR = os.path.join(ORIGINAL_DATASET_DIR, 'Normal')  
DATASET_DIR = os.path.join(BASE_PATH, 'small_dataset')  
TRAIN_DIR = os.path.join(DATASET_DIR, 'train')  
VALIDATION_DIR = os.path.join(DATASET_DIR, 'validation')  
TEST_DIR = os.path.join(DATASET_DIR, 'test')  
TRAIN_VIRAL_DIR = os.path.join(TRAIN_DIR, 'viral_pneumonia')  
TRAIN_COVID_DIR = os.path.join(TRAIN_DIR, 'covid')  
TRAIN_NORMAL_DIR = os.path.join(TRAIN_DIR, 'normal')  
VALIDATION_VIRAL_DIR = os.path.join(VALIDATION_DIR, 'viral_pneumonia')  
VALIDATION_COVID_DIR = os.path.join(VALIDATION_DIR, 'covid')  
VALIDATION_NORMAL_DIR = os.path.join(VALIDATION_DIR, 'normal')  
TEST_VIRAL_DIR = os.path.join(TEST_DIR, 'viral_pneumonia')  
TEST_COVID_DIR = os.path.join(TEST_DIR, 'covid')  
TEST_NORMAL_DIR = os.path.join(TEST_DIR, 'normal')  
  
def createDir(path: str) -> None:  
    if not os.path.exists(path):  
        os.mkdir(path)
```

```

createDir(DATASET_DIR)
createDir(TRAIN_DIR)
createDir(VALIDATION_DIR)
createDir(TEST_DIR)
createDir(TRAIN_VIRAL_DIR)
createDir(TRAIN_COVID_DIR)
createDir(TRAIN_NORMAL_DIR)
createDir(VALIDATION_VIRAL_DIR)
createDir(VALIDATION_COVID_DIR)
createDir(VALIDATION_NORMAL_DIR)
createDir(TEST_VIRAL_DIR)
createDir(TEST_COVID_DIR)
createDir(TEST_NORMAL_DIR)

```

```

[ ]: import numpy as np
import shutil

def generate_sets(source: str):
    allFiles = os.listdir(source)
    np.random.shuffle(allFiles)
    return np.split(np.array(allFiles), [int(len(allFiles)*0.7),
    ↪int(len(allFiles)*0.85)])

def saveAndSeparateFiles(src_dir: str, train_dir: str, val_dir: str, test_dir):
    train_fnames, val_fnames, test_fnames = generate_sets(src_dir)
    for fname in train_fnames:
        src = os.path.join(src_dir, fname)
        dst = os.path.join(train_dir, fname)
        shutil.copyfile(src, dst)

    for fname in val_fnames:
        src = os.path.join(src_dir, fname)
        dst = os.path.join(val_dir, fname)
        shutil.copyfile(src, dst)

    for fname in test_fnames:
        src = os.path.join(src_dir, fname)
        dst = os.path.join(test_dir, fname)
        shutil.copyfile(src, dst)

create = False
if create:
    saveAndSeparateFiles(ORIGINAL_NORMAL_DIR, TRAIN_NORMAL_DIR,
                        VALIDATION_NORMAL_DIR, TEST_NORMAL_DIR)

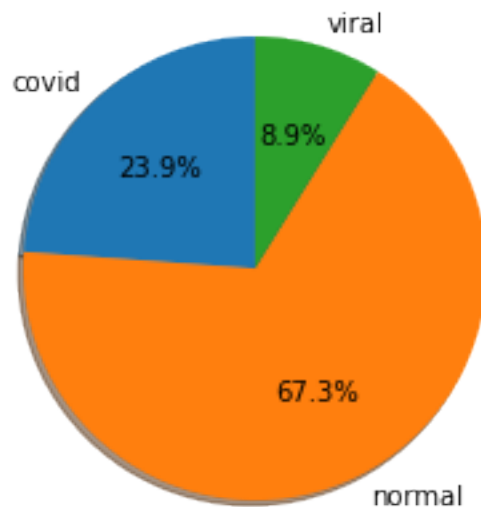
```

```
saveAndSeparateFiles(ORIGINAL_COVID_DIR, TRAIN_COVID_DIR,  
                     VALIDATION_COVID_DIR, TEST_COVID_DIR)  
saveAndSeparateFiles(ORIGINAL_VIRAL_DIR, TRAIN_VIRAL_DIR,  
                     VALIDATION_VIRAL_DIR, TEST_VIRAL_DIR)
```

1.2 Counting our images

```
[ ]: import tensorflow as tf  
import matplotlib.pyplot as plt  
normal_train = tf.io.gfile.glob(TRAIN_NORMAL_DIR + '/*')  
viral_train = tf.io.gfile.glob(TRAIN_VIRAL_DIR + '/*')  
covid_train = tf.io.gfile.glob(TRAIN_COVID_DIR + '/*')  
  
# Plotting Distribution of Each Classes  
image_count = {'covid': len(covid_train), 'normal': len(  
    normal_train), 'viral': len(viral_train)}  
fig1, ax1 = plt.subplots()  
ax1.pie(image_count.values(),  
        labels=image_count.keys(),  
        shadow=True,  
        autopct='%1.1f%%',  
        startangle=90)  
plt.show()
```

2021-10-07 20:14:24.777820: I
tensorflow/stream_executor/platform/default/dso_loader.cc:49] Successfully
opened dynamic library libcudart.so.10.1



1.3 Create our Covnet Model

In this case we are doing a multi class classification, our total classes are 3: - Viral CXR - Covid CXR - Normal CXR

Our neural network will output neurons as 3 classes that will calculate the probability of being one using the softmax function.

```
[ ]: from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(
    rescale=1./255,
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    # randomly rotate images in the range (degrees, 0 to 180)
    rotation_range=10,
    zoom_range=0.1, # Randomly zoom image
    # randomly shift images horizontally (fraction of total width)
    width_shift_range=0.1,
    # randomly shift images vertically (fraction of total height)
    height_shift_range=0.1,
    horizontal_flip=False, # randomly flip images
    vertical_flip=False # randomly flip images
)

# train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)
evaluate_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    TRAIN_DIR,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    color_mode='grayscale'
)

print(train_generator.class_indices)

validation_generator = test_datagen.flow_from_directory(
    VALIDATION_DIR,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    color_mode='grayscale'
```

```

)

print(validation_generator.class_indices)

test_generator = evaluate_datagen.flow_from_directory(
    TEST_DIR,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    color_mode='grayscale'
)

print(test_generator.class_indices)


```

Found 10606 images belonging to 3 classes.
{'covid': 0, 'normal': 1, 'viral_pneumonia': 2}
Found 2273 images belonging to 3 classes.
{'covid': 0, 'normal': 1, 'viral_pneumonia': 2}
Found 2274 images belonging to 3 classes.
{'covid': 0, 'normal': 1, 'viral_pneumonia': 2}

```

[ ]: for X_batch, y_batch in train_generator:
        # create a grid of 3x3 images
        for i in range(0, 9):
            plt.subplot(330 + 1 + i)
            plt.imshow(X_batch[i].reshape(150, 150), cmap=plt.
↪get_cmap('gray'))
            # show the plot
            plt.show()
        break

```



```
[ ]: from keras.layers import Conv2D, BatchNormalization, MaxPooling2D, Dropout, Flatten, Dense
    from keras.models import Sequential
    from keras import backend

    backend.clear_session()

    model = Sequential()
    model.add(Conv2D(64, (3, 3), activation='relu', input_shape=(150, 150, 1)))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Conv2D(64, (3, 3), activation='relu'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Conv2D(128, (3, 3), activation='relu'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Conv2D(128, (3, 3), activation='relu'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Flatten())
    model.add(Dropout(0.5))
    model.add(Dense(512, activation='relu'))
    model.add(Dense(64, activation='relu'))
    model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|---|----------------------|---------|
| conv2d (Conv2D) | (None, 148, 148, 64) | 640 |
| batch_normalization (Batch Normalization) | (None, 148, 148, 64) | 256 |
| max_pooling2d (MaxPooling2D) | (None, 74, 74, 64) | 0 |
| conv2d_1 (Conv2D) | (None, 72, 72, 64) | 36928 |
| batch_normalization_1 (Batch Normalization) | (None, 72, 72, 64) | 256 |
| max_pooling2d_1 (MaxPooling2D) | (None, 36, 36, 64) | 0 |
| conv2d_2 (Conv2D) | (None, 34, 34, 128) | 73856 |
| batch_normalization_2 (Batch Normalization) | (None, 34, 34, 128) | 512 |
| max_pooling2d_2 (MaxPooling2D) | (None, 17, 17, 128) | 0 |
| conv2d_3 (Conv2D) | (None, 15, 15, 128) | 147584 |
| batch_normalization_3 (Batch Normalization) | (None, 15, 15, 128) | 512 |
| max_pooling2d_3 (MaxPooling2D) | (None, 7, 7, 128) | 0 |
| flatten (Flatten) | (None, 6272) | 0 |
| dropout (Dropout) | (None, 6272) | 0 |
| dense (Dense) | (None, 512) | 3211776 |
| dense_1 (Dense) | (None, 64) | 32832 |
| dense_2 (Dense) | (None, 3) | 195 |

Total params: 3,505,347
Trainable params: 3,504,579
Non-trainable params: 768

```
[ ]: from keras import optimizers
```

```
# opt = RMSprop(lr=0.0001, decay=1e-6)
opt = optimizers.Adam(learning_rate=1e-5, decay=1e-7)

model.compile(loss='categorical_crossentropy', optimizer=opt,
↳metrics=['accuracy'])
```

```
[ ]: import numpy as np
from sklearn.utils import class_weight
from keras.callbacks import EarlyStopping
from keras.callbacks import ModelCheckpoint

classes = train_generator.classes
class_weights = class_weight.compute_class_weight(None,
                                                    np.unique(classes),
                                                    classes)

best_model_path = os.path.join(BASE_PATH, 'best_model.h5')
es = EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=8)
mc = ModelCheckpoint(best_model_path, monitor='val_accuracy', mode='max',
↳verbose=1, save_best_only=True)
history = model.fit(
    train_generator,
    steps_per_epoch=train_generator.n // 32,
    epochs=150,
    validation_data=validation_generator,
    class_weight=dict(zip(np.unique(classes), class_weights)),
    callbacks=[es, mc]
)
```

Epoch 1/150

331/331 [=====] - 36s 108ms/step - loss: 1.1811 -
accuracy: 0.5959 - val_loss: 1.6331 - val_accuracy: 0.2411

Epoch 00001: val_accuracy improved from -inf to 0.24109, saving model to
/home/hivini/learn/research/new-covid/best_model.h5

Epoch 2/150

331/331 [=====] - 33s 101ms/step - loss: 0.6353 -
accuracy: 0.7354 - val_loss: 0.4355 - val_accuracy: 0.8170

Epoch 00002: val_accuracy improved from 0.24109 to 0.81698, saving model to
/home/hivini/learn/research/new-covid/best_model.h5

Epoch 3/150

331/331 [=====] - 32s 97ms/step - loss: 0.5500 -
accuracy: 0.7627 - val_loss: 0.3936 - val_accuracy: 0.8214

Epoch 00003: val_accuracy improved from 0.81698 to 0.82138, saving model to
/home/hivini/learn/research/new-covid/best_model.h5

Epoch 4/150

331/331 [=====] - 32s 97ms/step - loss: 0.4964 - accuracy: 0.7868 - val_loss: 0.3578 - val_accuracy: 0.8460

Epoch 00004: val_accuracy improved from 0.82138 to 0.84602, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 5/150

331/331 [=====] - 33s 98ms/step - loss: 0.4639 - accuracy: 0.8095 - val_loss: 0.3470 - val_accuracy: 0.8482

Epoch 00005: val_accuracy improved from 0.84602 to 0.84822, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 6/150

331/331 [=====] - 32s 97ms/step - loss: 0.4297 - accuracy: 0.8165 - val_loss: 0.3197 - val_accuracy: 0.8614

Epoch 00006: val_accuracy improved from 0.84822 to 0.86142, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 7/150

331/331 [=====] - 32s 98ms/step - loss: 0.4217 - accuracy: 0.8239 - val_loss: 0.3027 - val_accuracy: 0.8667

Epoch 00007: val_accuracy improved from 0.86142 to 0.86670, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 8/150

331/331 [=====] - 32s 98ms/step - loss: 0.3966 - accuracy: 0.8354 - val_loss: 0.2895 - val_accuracy: 0.8759

Epoch 00008: val_accuracy improved from 0.86670 to 0.87593, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 9/150

331/331 [=====] - 32s 97ms/step - loss: 0.3773 - accuracy: 0.8380 - val_loss: 0.2666 - val_accuracy: 0.8852

Epoch 00009: val_accuracy improved from 0.87593 to 0.88517, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 10/150

331/331 [=====] - 32s 97ms/step - loss: 0.3625 - accuracy: 0.8536 - val_loss: 0.2635 - val_accuracy: 0.8861

Epoch 00010: val_accuracy improved from 0.88517 to 0.88605, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 11/150

331/331 [=====] - 32s 98ms/step - loss: 0.3531 - accuracy: 0.8498 - val_loss: 0.2518 - val_accuracy: 0.8962

Epoch 00011: val_accuracy improved from 0.88605 to 0.89617, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 12/150

331/331 [=====] - 32s 98ms/step - loss: 0.3289 - accuracy: 0.8685 - val_loss: 0.2681 - val_accuracy: 0.8755

Epoch 00012: val_accuracy did not improve from 0.89617
Epoch 13/150

331/331 [=====] - 33s 99ms/step - loss: 0.3219 - accuracy: 0.8657 - val_loss: 0.2387 - val_accuracy: 0.9001

Epoch 00013: val_accuracy improved from 0.89617 to 0.90013, saving model to /home/hivini/learn/research/new-covid/best_model.h5
Epoch 14/150

331/331 [=====] - 32s 98ms/step - loss: 0.3104 - accuracy: 0.8723 - val_loss: 0.2520 - val_accuracy: 0.8869

Epoch 00014: val_accuracy did not improve from 0.90013
Epoch 15/150

331/331 [=====] - 32s 98ms/step - loss: 0.2965 - accuracy: 0.8830 - val_loss: 0.2434 - val_accuracy: 0.8971

Epoch 00015: val_accuracy did not improve from 0.90013
Epoch 16/150

331/331 [=====] - 34s 101ms/step - loss: 0.2969 - accuracy: 0.8806 - val_loss: 0.2118 - val_accuracy: 0.9147

Epoch 00016: val_accuracy improved from 0.90013 to 0.91465, saving model to /home/hivini/learn/research/new-covid/best_model.h5
Epoch 17/150

331/331 [=====] - 32s 97ms/step - loss: 0.2868 - accuracy: 0.8831 - val_loss: 0.2367 - val_accuracy: 0.9028

Epoch 00017: val_accuracy did not improve from 0.91465
Epoch 18/150

331/331 [=====] - 32s 97ms/step - loss: 0.2923 - accuracy: 0.8842 - val_loss: 0.2188 - val_accuracy: 0.9138

Epoch 00018: val_accuracy did not improve from 0.91465
Epoch 19/150

331/331 [=====] - 32s 97ms/step - loss: 0.2632 - accuracy: 0.8945 - val_loss: 0.2025 - val_accuracy: 0.9243

Epoch 00019: val_accuracy improved from 0.91465 to 0.92433, saving model to /home/hivini/learn/research/new-covid/best_model.h5
Epoch 20/150

331/331 [=====] - 32s 97ms/step - loss: 0.2504 - accuracy: 0.8963 - val_loss: 0.2460 - val_accuracy: 0.9001

Epoch 00020: val_accuracy did not improve from 0.92433
Epoch 21/150

331/331 [=====] - 33s 100ms/step - loss: 0.2438 -
accuracy: 0.9013 - val_loss: 0.2119 - val_accuracy: 0.9173

Epoch 00021: val_accuracy did not improve from 0.92433

Epoch 22/150

331/331 [=====] - 39s 117ms/step - loss: 0.2328 -
accuracy: 0.9100 - val_loss: 0.2993 - val_accuracy: 0.8869

Epoch 00022: val_accuracy did not improve from 0.92433

Epoch 23/150

331/331 [=====] - 33s 101ms/step - loss: 0.2528 -
accuracy: 0.9013 - val_loss: 0.2243 - val_accuracy: 0.9081

Epoch 00023: val_accuracy did not improve from 0.92433

Epoch 24/150

331/331 [=====] - 32s 97ms/step - loss: 0.2335 -
accuracy: 0.9075 - val_loss: 0.1725 - val_accuracy: 0.9344

Epoch 00024: val_accuracy improved from 0.92433 to 0.93445, saving model to
/home/hivini/learn/research/new-covid/best_model.h5

Epoch 25/150

331/331 [=====] - 32s 97ms/step - loss: 0.2201 -
accuracy: 0.9155 - val_loss: 0.2458 - val_accuracy: 0.9045

Epoch 00025: val_accuracy did not improve from 0.93445

Epoch 26/150

331/331 [=====] - 32s 97ms/step - loss: 0.2123 -
accuracy: 0.9144 - val_loss: 0.1720 - val_accuracy: 0.9340

Epoch 00026: val_accuracy did not improve from 0.93445

Epoch 27/150

331/331 [=====] - 32s 97ms/step - loss: 0.2226 -
accuracy: 0.9140 - val_loss: 0.2461 - val_accuracy: 0.9028

Epoch 00027: val_accuracy did not improve from 0.93445

Epoch 28/150

331/331 [=====] - 32s 97ms/step - loss: 0.2106 -
accuracy: 0.9187 - val_loss: 0.1939 - val_accuracy: 0.9278

Epoch 00028: val_accuracy did not improve from 0.93445

Epoch 29/150

331/331 [=====] - 32s 97ms/step - loss: 0.2133 -
accuracy: 0.9168 - val_loss: 0.1702 - val_accuracy: 0.9362

Epoch 00029: val_accuracy improved from 0.93445 to 0.93621, saving model to
/home/hivini/learn/research/new-covid/best_model.h5

Epoch 30/150

331/331 [=====] - 32s 97ms/step - loss: 0.2060 -

accuracy: 0.9200 - val_loss: 0.2272 - val_accuracy: 0.9168

Epoch 00030: val_accuracy did not improve from 0.93621

Epoch 31/150

331/331 [=====] - 33s 100ms/step - loss: 0.1989 -
accuracy: 0.9221 - val_loss: 0.1639 - val_accuracy: 0.9419

Epoch 00031: val_accuracy improved from 0.93621 to 0.94193, saving model to
/home/hivini/learn/research/new-covid/best_model.h5

Epoch 32/150

331/331 [=====] - 32s 98ms/step - loss: 0.2006 -
accuracy: 0.9201 - val_loss: 0.1861 - val_accuracy: 0.9296

Epoch 00032: val_accuracy did not improve from 0.94193

Epoch 33/150

331/331 [=====] - 32s 98ms/step - loss: 0.1900 -
accuracy: 0.9280 - val_loss: 0.1598 - val_accuracy: 0.9393

Epoch 00033: val_accuracy did not improve from 0.94193

Epoch 34/150

331/331 [=====] - 32s 98ms/step - loss: 0.1854 -
accuracy: 0.9290 - val_loss: 0.1560 - val_accuracy: 0.9415

Epoch 00034: val_accuracy did not improve from 0.94193

Epoch 35/150

331/331 [=====] - 32s 97ms/step - loss: 0.1928 -
accuracy: 0.9256 - val_loss: 0.1750 - val_accuracy: 0.9362

Epoch 00035: val_accuracy did not improve from 0.94193

Epoch 36/150

331/331 [=====] - 32s 97ms/step - loss: 0.1756 -
accuracy: 0.9320 - val_loss: 0.1888 - val_accuracy: 0.9300

Epoch 00036: val_accuracy did not improve from 0.94193

Epoch 37/150

331/331 [=====] - 32s 98ms/step - loss: 0.1704 -
accuracy: 0.9343 - val_loss: 0.1431 - val_accuracy: 0.9485

Epoch 00037: val_accuracy improved from 0.94193 to 0.94853, saving model to
/home/hivini/learn/research/new-covid/best_model.h5

Epoch 38/150

331/331 [=====] - 32s 97ms/step - loss: 0.1706 -
accuracy: 0.9365 - val_loss: 0.2093 - val_accuracy: 0.9182

Epoch 00038: val_accuracy did not improve from 0.94853

Epoch 39/150

331/331 [=====] - 33s 99ms/step - loss: 0.1698 -
accuracy: 0.9343 - val_loss: 0.1483 - val_accuracy: 0.9459

Epoch 00039: val_accuracy did not improve from 0.94853
Epoch 40/150
331/331 [=====] - 33s 99ms/step - loss: 0.1831 -
accuracy: 0.9300 - val_loss: 0.3303 - val_accuracy: 0.8808

Epoch 00040: val_accuracy did not improve from 0.94853
Epoch 41/150
331/331 [=====] - 32s 98ms/step - loss: 0.1693 -
accuracy: 0.9342 - val_loss: 0.1278 - val_accuracy: 0.9538

Epoch 00041: val_accuracy improved from 0.94853 to 0.95381, saving model to
/home/hivini/learn/research/new-covid/best_model.h5
Epoch 42/150
331/331 [=====] - 32s 98ms/step - loss: 0.1604 -
accuracy: 0.9406 - val_loss: 0.1282 - val_accuracy: 0.9529

Epoch 00042: val_accuracy did not improve from 0.95381
Epoch 43/150
331/331 [=====] - 32s 98ms/step - loss: 0.1759 -
accuracy: 0.9362 - val_loss: 0.1600 - val_accuracy: 0.9459

Epoch 00043: val_accuracy did not improve from 0.95381
Epoch 44/150
331/331 [=====] - 32s 98ms/step - loss: 0.1678 -
accuracy: 0.9350 - val_loss: 0.1712 - val_accuracy: 0.9388

Epoch 00044: val_accuracy did not improve from 0.95381
Epoch 45/150
331/331 [=====] - 32s 98ms/step - loss: 0.1472 -
accuracy: 0.9466 - val_loss: 0.2125 - val_accuracy: 0.9221

Epoch 00045: val_accuracy did not improve from 0.95381
Epoch 46/150
331/331 [=====] - 34s 104ms/step - loss: 0.1656 -
accuracy: 0.9405 - val_loss: 0.1405 - val_accuracy: 0.9463

Epoch 00046: val_accuracy did not improve from 0.95381
Epoch 47/150
331/331 [=====] - 38s 114ms/step - loss: 0.1587 -
accuracy: 0.9422 - val_loss: 0.1406 - val_accuracy: 0.9454

Epoch 00047: val_accuracy did not improve from 0.95381
Epoch 48/150
331/331 [=====] - 32s 97ms/step - loss: 0.1546 -
accuracy: 0.9406 - val_loss: 0.1440 - val_accuracy: 0.9415

Epoch 00048: val_accuracy did not improve from 0.95381

Epoch 49/150
 331/331 [=====] - 33s 98ms/step - loss: 0.1411 - accuracy: 0.9506 - val_loss: 0.1244 - val_accuracy: 0.9520

Epoch 00049: val_accuracy did not improve from 0.95381

Epoch 50/150
 331/331 [=====] - 32s 98ms/step - loss: 0.1466 - accuracy: 0.9471 - val_loss: 0.1227 - val_accuracy: 0.9542

Epoch 00050: val_accuracy improved from 0.95381 to 0.95425, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 51/150
 331/331 [=====] - 32s 98ms/step - loss: 0.1399 - accuracy: 0.9502 - val_loss: 0.1145 - val_accuracy: 0.9604

Epoch 00051: val_accuracy improved from 0.95425 to 0.96040, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 52/150
 331/331 [=====] - 33s 98ms/step - loss: 0.1493 - accuracy: 0.9454 - val_loss: 0.1282 - val_accuracy: 0.9534

Epoch 00052: val_accuracy did not improve from 0.96040

Epoch 53/150
 331/331 [=====] - 32s 97ms/step - loss: 0.1451 - accuracy: 0.9471 - val_loss: 0.1304 - val_accuracy: 0.9525

Epoch 00053: val_accuracy did not improve from 0.96040

Epoch 54/150
 331/331 [=====] - 32s 97ms/step - loss: 0.1453 - accuracy: 0.9455 - val_loss: 0.1264 - val_accuracy: 0.9534

Epoch 00054: val_accuracy did not improve from 0.96040

Epoch 55/150
 331/331 [=====] - 33s 99ms/step - loss: 0.1311 - accuracy: 0.9511 - val_loss: 0.1231 - val_accuracy: 0.9556

Epoch 00055: val_accuracy did not improve from 0.96040

Epoch 56/150
 331/331 [=====] - 33s 100ms/step - loss: 0.1331 - accuracy: 0.9511 - val_loss: 0.1078 - val_accuracy: 0.9630

Epoch 00056: val_accuracy improved from 0.96040 to 0.96304, saving model to /home/hivini/learn/research/new-covid/best_model.h5

Epoch 57/150
 331/331 [=====] - 33s 100ms/step - loss: 0.1352 - accuracy: 0.9520 - val_loss: 0.1473 - val_accuracy: 0.9424

Epoch 00057: val_accuracy did not improve from 0.96304

Epoch 58/150
331/331 [=====] - 34s 102ms/step - loss: 0.1304 -
accuracy: 0.9513 - val_loss: 0.3040 - val_accuracy: 0.8883

Epoch 00058: val_accuracy did not improve from 0.96304
Epoch 59/150
331/331 [=====] - 33s 100ms/step - loss: 0.1319 -
accuracy: 0.9511 - val_loss: 0.1260 - val_accuracy: 0.9551

Epoch 00059: val_accuracy did not improve from 0.96304
Epoch 60/150
331/331 [=====] - 33s 99ms/step - loss: 0.1343 -
accuracy: 0.9503 - val_loss: 0.2046 - val_accuracy: 0.9212

Epoch 00060: val_accuracy did not improve from 0.96304
Epoch 61/150
331/331 [=====] - 33s 100ms/step - loss: 0.1285 -
accuracy: 0.9499 - val_loss: 0.1233 - val_accuracy: 0.9534

Epoch 00061: val_accuracy did not improve from 0.96304
Epoch 62/150
331/331 [=====] - 33s 98ms/step - loss: 0.1302 -
accuracy: 0.9519 - val_loss: 0.1722 - val_accuracy: 0.9305

Epoch 00062: val_accuracy did not improve from 0.96304
Epoch 63/150
331/331 [=====] - 32s 98ms/step - loss: 0.1256 -
accuracy: 0.9552 - val_loss: 0.1041 - val_accuracy: 0.9626

Epoch 00063: val_accuracy did not improve from 0.96304
Epoch 64/150
331/331 [=====] - 33s 98ms/step - loss: 0.1235 -
accuracy: 0.9542 - val_loss: 0.1200 - val_accuracy: 0.9573

Epoch 00064: val_accuracy did not improve from 0.96304
Epoch 65/150
331/331 [=====] - 32s 97ms/step - loss: 0.1157 -
accuracy: 0.9580 - val_loss: 0.1047 - val_accuracy: 0.9604

Epoch 00065: val_accuracy did not improve from 0.96304
Epoch 66/150
331/331 [=====] - 32s 97ms/step - loss: 0.1067 -
accuracy: 0.9613 - val_loss: 0.1170 - val_accuracy: 0.9564

Epoch 00066: val_accuracy did not improve from 0.96304
Epoch 67/150
331/331 [=====] - 32s 97ms/step - loss: 0.1211 -
accuracy: 0.9561 - val_loss: 0.1134 - val_accuracy: 0.9613

Epoch 00067: val_accuracy did not improve from 0.96304

Epoch 68/150

331/331 [=====] - 32s 97ms/step - loss: 0.1259 - accuracy: 0.9533 - val_loss: 0.1126 - val_accuracy: 0.9591

Epoch 00068: val_accuracy did not improve from 0.96304

Epoch 69/150

331/331 [=====] - 33s 100ms/step - loss: 0.1228 - accuracy: 0.9567 - val_loss: 0.1229 - val_accuracy: 0.9542

Epoch 00069: val_accuracy did not improve from 0.96304

Epoch 70/150

331/331 [=====] - 34s 101ms/step - loss: 0.1139 - accuracy: 0.9597 - val_loss: 0.1443 - val_accuracy: 0.9450

Epoch 00070: val_accuracy did not improve from 0.96304

Epoch 71/150

331/331 [=====] - 33s 100ms/step - loss: 0.1144 - accuracy: 0.9566 - val_loss: 0.1493 - val_accuracy: 0.9468

Epoch 00071: val_accuracy did not improve from 0.96304

Epoch 00071: early stopping

```
[ ]: model.save(os.path.join(BASE_PATH, 'covid_classifier_result.h5'))
```

```
[ ]: test_loss, test_acc = model.evaluate(test_generator)
print("Loss on test set: ", test_loss)
print("Accuracy on test set: ", test_acc)
```

72/72 [=====] - 4s 51ms/step - loss: 0.1349 - accuracy: 0.9529

Loss on test set: 0.1348775327205658

Accuracy on test set: 0.9529463648796082

```
[ ]: import matplotlib.pyplot as plt


acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)
# bo is for blue dot.
plt.plot(epochs, loss, 'g', label='Training loss')
# b is for solid blue line
plt.plot(epochs, val_loss, 'b', label='Validation loss')
```



```
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()


plt.show()
```



```
[ ]: plt.clf()

plt.plot(epochs, acc, 'g', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



```
[ ]: from sklearn.metrics import classification_report, confusion_matrix

X_test = []
Y_test = []
# Extract the data
for X, Y in test_generator:
    X_test.append(X)
    Y_test.append(Y)
X_test = np.array(X_test)
Y_test = np.array(Y_test)

predictions = model.predict_classes(X_test)
predictions = predictions.reshape(1, -1)[0]

print(classification_report(Y_test, predictions, target_names=[
    'Covid (Class 0)', 'Normal (Class 1)', 'Viral Pneumonia (Class 2)']))
```

```
[ ]: import pandas as pd
import seaborn as sns

labels = ['covid', 'normal', 'viral_pneumonia']

cm = confusion_matrix(Y_test, predictions)
```

```
cm = pd.DataFrame(cm , index = ['0','1'] , columns = ['0','1'])
plt.figure(figsize = (10,10))
sns.heatmap(cm,cmap= "Reds", linecolor = 'black' , linewidth = 1 , annot = ␣
↪True, fmt='',xticklabels = labels,yticklabels = labels)
```